

▼ Image Classification

```
import matplotlib.pyplot as plt
import numpy as np
import PIL
import tensorflow as tf
import pathlib

from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.models import Sequential
```

▼ Explore Dataset

This dataset contains images of flowers, split into 5 classes. The expected prediction is which species of flower is pictured.

```
dataset_url = "https://storage.googleapis.com/download.tensorflow.org/example_data/flower_photos.tar.gz"
data_dir = tf.keras.utils.get_file('flower_photos.tar', origin=dataset_url,
data_dir = pathlib.Path(data_dir).with_suffix(''))

image_count = len(list(data_dir.glob('*/*.jpg')))
print(image_count)
```

```
Downloading data from https://storage.googleapis.com/download.tensorflow.org/example\_data/flower\_photos.tar.gz
228813984/228813984 [=====] - 2s 0us/step
3670
```

```
batch_size = 32
img_height = 180
img_width = 180
```

```
train_ds = tf.keras.utils.image_dataset_from_directory(
    data_dir,
    validation_split=0.2,
    subset="training",
    seed=123,
    image_size=(img_height, img_width),
    batch_size=batch_size)
```

```
Found 3670 files belonging to 5 classes.
Using 2936 files for training.
```

✓ 3s completed at 5:02 PM



```
val_ds = tf.keras.utils.image_dataset_from_directory(
    data_dir,
    validation_split=0.2,
    subset="validation",
    seed=123,
    image_size=(img_height, img_width),
    batch_size=batch_size)
```

Found 3670 files belonging to 5 classes.
Using 734 files for validation.

```
class_names = train_ds.class_names
print(class_names)
```

```
['daisy', 'dandelion', 'roses', 'sunflowers', 'tulips']
```

```
plt.figure(figsize=(10, 10))
for images, labels in train_ds.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i].numpy().astype("uint8"))
        plt.title(class_names[labels[i]])
        plt.axis("off")
```

roses



dandelion



tulips



sunflowers



dandelion



roses



dandelion

roses

tulips



Tune and Run Model

```
AUTOTUNE = tf.data.AUTOTUNE
```

```
train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)  
val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
```

```
normalization_layer = layers.Rescaling(1./255)
```

```
normalized_ds = train_ds.map(lambda x, y: (normalization_layer(x), y))  
image_batch, labels_batch = next(iter(normalized_ds))  
first_image = image_batch[0]  
print(np.min(first_image), np.max(first_image))
```

```
0.0 0.9993303
```

```
num_classes = len(class_names)
```

```
model = Sequential([  
    layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),  
    layers.Conv2D(16, 3, padding='same', activation='relu'),  
    layers.MaxPooling2D(),  
    layers.Conv2D(32, 3, padding='same', activation='relu'),  
    layers.MaxPooling2D(),  
    layers.Conv2D(64, 3, padding='same', activation='relu'),  
    layers.MaxPooling2D(),  
    layers.Flatten(),
```

```
layers.Dense(128, activation='relu'),  
layers.Dense(num_classes)  
)
```

```
model.compile(optimizer='adam',  
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),  
              metrics=['accuracy'])
```

```
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
rescaling_3 (Rescaling)	(None, 180, 180, 3)	0
conv2d_3 (Conv2D)	(None, 180, 180, 16)	448
max_pooling2d_3 (MaxPooling2D)	(None, 90, 90, 16)	0
conv2d_4 (Conv2D)	(None, 90, 90, 32)	4640
max_pooling2d_4 (MaxPooling2D)	(None, 45, 45, 32)	0
conv2d_5 (Conv2D)	(None, 45, 45, 64)	18496
max_pooling2d_5 (MaxPooling2D)	(None, 22, 22, 64)	0
flatten_1 (Flatten)	(None, 30976)	0
dense_2 (Dense)	(None, 128)	3965056
dense_3 (Dense)	(None, 5)	645

```
Total params: 3,989,285  
Trainable params: 3,989,285  
Non-trainable params: 0
```

```
epochs=10  
history = model.fit(  
    train_ds,  
    validation_data=val_ds,  
    epochs=epochs  
)
```

```
Epoch 8/10  
92/92 [=====] - 105s 1s/step - loss: 0.0893 -  
Epoch 9/10  
92/92 [=====] - 105s 1s/step - loss: 0.0358 -  
Epoch 10/10  
92/92 [=====] - 110s 1s/step - loss: 0.0225 -
```


```
acc = history.history['accuracy']  
val_acc = history.history['val_accuracy']
```


```
loss = history.history['loss']  
val_loss = history.history['val_loss']
```

```
epochs_range = range(epochs)
```

```
plt.figure(figsize=(8, 8))  
plt.subplot(1, 2, 1)  
plt.plot(epochs_range, acc, label='Training Accuracy')  
plt.plot(epochs_range, val_acc, label='Validation Accuracy')  
plt.legend(loc='lower right')  
plt.title('Training and Validation Accuracy')
```

```
plt.subplot(1, 2, 2)  
plt.plot(epochs_range, loss, label='Training Loss')  
plt.plot(epochs_range, val_loss, label='Validation Loss')  
plt.legend(loc='upper right')  
plt.title('Training and Validation Loss')  
plt.show()
```





The accuracy of this model is 60%. This can be improved through further training and data augmentation.

The accuracy of this model is 60%. This can be improved through further training and data augmentation.

Transfer Learning

```
base_model = tf.keras.applications.MobileNetV2(input_shape=(img_height, img_width, 3),
                                                include_top=False,
                                                weights='imagenet')

base_model.trainable = False
base_model.summary()
```

WARNING:tensorflow: `input_shape` is undefined or non-square, or `rows` is not divisible by the image's height. Please specify a valid `input_shape`. See also https://www.tensorflow.org/api_guides/python/preprocessing_image.

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet/mobilenetv2_1.00_224.tgz [9406464/9406464] - 0s 0us/step

Model: "mobilenetv2_1.00_224"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 180, 180, 3)]	0	None
Conv1 (Conv2D)	(None, 90, 90, 32)	864	input_1
bn_Conv1 (BatchNormalization)	(None, 90, 90, 32)	128	Conv1
Conv1_relu (ReLU)	(None, 90, 90, 32)	0	bn_Conv1

expanded_conv_depthwise (DepthwiseConv2D)	(None, 90, 90, 32)	288	['Conv
expanded_conv_depthwise_BN (BatchNormalization)	(None, 90, 90, 32)	128	['expa
expanded_conv_depthwise_relu (ReLU)	(None, 90, 90, 32)	0	['expa ']
expanded_conv_project (Conv2D)	(None, 90, 90, 16)	512	['expa [0]']
expanded_conv_project_BN (BatchNormalization)	(None, 90, 90, 16)	64	['expa
block_1_expand (Conv2D)	(None, 90, 90, 96)	1536	['expa ']
block_1_expand_BN (BatchNormalization)	(None, 90, 90, 96)	384	['bloc
block_1_expand_relu (ReLU)	(None, 90, 90, 96)	0	['bloc
block_1_pad (ZeroPadding2D)	(None, 91, 91, 96)	0	['bloc
block_1_depthwise (DepthwiseConv2D)	(None, 45, 45, 96)	864	['bloc
block_1_depthwise_BN (BatchNormalization)	(None, 45, 45, 96)	384	['bloc
block_1_depthwise_relu (ReLU)	(None, 45, 45, 96)	0	['bloc
block_1_project (Conv2D)	(None, 45, 45, 24)	2304	['bloc
block_1_project_BN (BatchNormalization)	(None, 45, 45, 24)	96	['bloc
block_2_expand (Conv2D)	(None, 45, 45, 144)	2456	['bloc

[Colab paid products](#) - [Cancel contracts here](#)